

0.PreparingDataOceanProximityClassification

This California Housing Prices dataset has been downloaded from StatLib repository (<http://lib.stat.cmu.edu/datasets/>). It is based on data from the 1990 California census. It is not recent, but this is not important for deep learning. The original dataset appeared in R. Kelley Pace and Ronald Barry, "Sparse Spatial Autoregressions," Statistics & Probability Letters 33, no. 3 (1997): 291–297.

Data for each instance (observation) is referred to a block group in California, which could be corresponded to a district, with a population of 600 to 3,000 people, and 1,425.5 on average.

OceanProximityPreparedCleanAttributes.csv The original dataset contained 20,640 instances, which is cleaned, preprocessed and prepared in this notebook. After this phase of data preparation, a final dataset of 20,433 instances are obtained with 9 attributes individually normalized with a min-max scaling, $\frac{x-min}{max-min}$: *longitude* and *latitude* (location), *median age*, *total rooms*, *total bedrooms*, *population*, *households*, *median income* and *median house value*.

From this data, the classification problem consists on estimating the location (label *ocean proximity*), categorized into five classes: "<1H OCEAN", "INLAND" "NEAR BAY", "NEAR OCEAN" and "ISLAND". We will see that label "ISLAND" only has 5 instances. Therefore, they are removed from the dataset. The remaining classes are labelled from 0 ("<1H OCEAN") to 3 ("NEAR OCEAN"), and one-hot encoded in OceanProximityOneHotEncodedClasses.csv file for supervised training models.

```
In [1]: import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, minmax_scale
import matplotlib.pyplot as plt
```

```
In [2]: INPUT_FILE_NAME = "HousingRawDataset.csv"
ATT_FILE_NAME = "OceanProximityPreparedCleanAttributes.csv"
ONE_HOT_ENCODED_LABEL_FILE_NAME = "OceanProximityOneHotEncodedClasses.csv"
```

```
In [3]: dataset = pd.read_csv(INPUT_FILE_NAME)
```

```
In [4]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
longitude          20640 non-null float64
latitude           20640 non-null float64
housing_median_age  20640 non-null float64
```

```

total_rooms      20640 non-null float64
total_bedrooms   20433 non-null float64
population       20640 non-null float64
households       20640 non-null float64
median_income    20640 non-null float64
median_house_value 20640 non-null float64
ocean_proximity  20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

In [5]: dataset.iloc[:10]

```

Out[5]:   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0    -122.23    37.88           41.0           880.0           129.0
1    -122.22    37.86           21.0          7099.0          1106.0
2    -122.24    37.85           52.0          1467.0           190.0
3    -122.25    37.85           52.0          1274.0           235.0
4    -122.25    37.85           52.0          1627.0           280.0
5    -122.25    37.85           52.0           919.0           213.0
6    -122.25    37.84           52.0          2535.0           489.0
7    -122.25    37.84           52.0          3104.0           687.0
8    -122.26    37.84           42.0          2555.0           665.0
9    -122.25    37.84           52.0          3549.0           707.0

      population  households  median_income  median_house_value  ocean_proximity
0           322.0        126.0         8.3252         452600.0        NEAR BAY
1          2401.0       1138.0         8.3014         358500.0        NEAR BAY
2           496.0        177.0         7.2574         352100.0        NEAR BAY
3           558.0        219.0         5.6431         341300.0        NEAR BAY
4           565.0        259.0         3.8462         342200.0        NEAR BAY
5           413.0        193.0         4.0368         269700.0        NEAR BAY
6          1094.0        514.0         3.6591         299200.0        NEAR BAY
7          1157.0        647.0         3.1200         241400.0        NEAR BAY
8          1206.0        595.0         2.0804         226700.0        NEAR BAY
9          1551.0        714.0         3.6912         261100.0        NEAR BAY

```

In [6]: dataset.iloc[-10:]

```

Out[6]:   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
20630    -121.32    39.29           11.0           2640.0           505.0
20631    -121.40    39.33           15.0           2655.0           493.0
20632    -121.45    39.26           15.0           2319.0           416.0
20633    -121.53    39.19           27.0           2080.0           412.0
20634    -121.56    39.27           28.0           2332.0           395.0
20635    -121.09    39.48           25.0           1665.0           374.0
20636    -121.21    39.49           18.0            697.0           150.0
20637    -121.22    39.43           17.0           2254.0           485.0
20638    -121.32    39.43           18.0           1860.0           409.0

```

20639	-121.24	39.37	16.0	2785.0	616.0
-------	---------	-------	------	--------	-------

	population	households	median_income	median_house_value	\
20630	1257.0	445.0	3.5673	112000.0	
20631	1200.0	432.0	3.5179	107200.0	
20632	1047.0	385.0	3.1250	115600.0	
20633	1082.0	382.0	2.5495	98300.0	
20634	1041.0	344.0	3.7125	116800.0	
20635	845.0	330.0	1.5603	78100.0	
20636	356.0	114.0	2.5568	77100.0	
20637	1007.0	433.0	1.7000	92300.0	
20638	741.0	349.0	1.8672	84700.0	
20639	1387.0	530.0	2.3886	89400.0	

	ocean_proximity
20630	INLAND
20631	INLAND
20632	INLAND
20633	INLAND
20634	INLAND
20635	INLAND
20636	INLAND
20637	INLAND
20638	INLAND
20639	INLAND

In [7]: dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
longitude      20640 non-null float64
latitude       20640 non-null float64
housing_median_age  20640 non-null float64
total_rooms    20640 non-null float64
total_bedrooms 20433 non-null float64
population     20640 non-null float64
households     20640 non-null float64
median_income  20640 non-null float64
median_house_value 20640 non-null float64
ocean_proximity 20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

In [8]: labels=list(set(dataset["ocean_proximity"]))
labels

Out[8]: ['ISLAND', 'NEAR OCEAN', 'NEAR BAY', '<1H OCEAN', 'INLAND']

First Step: find out wheather or not there are missing values and, in such case, remove them

```
In [9]: {att : dataset[dataset[att].isnull()].shape[0] for att in dataset.columns}
```

```
Out[9]: {'longitude': 0,  
        'latitude': 0,  
        'housing_median_age': 0,  
        'total_rooms': 0,  
        'total_bedrooms': 207,  
        'population': 0,  
        'households': 0,  
        'median_income': 0,  
        'median_house_value': 0,  
        'ocean_proximity': 0}
```

total_bedrooms has 207 missing values. The corresponding rows are removed.

```
In [10]: dataset.dropna(inplace=True)
```

```
In [11]: {att : dataset[dataset[att].isnull()].shape[0] for att in dataset.columns}
```

```
Out[11]: {'longitude': 0,  
        'latitude': 0,  
        'housing_median_age': 0,  
        'total_rooms': 0,  
        'total_bedrooms': 0,  
        'population': 0,  
        'households': 0,  
        'median_income': 0,  
        'median_house_value': 0,  
        'ocean_proximity': 0}
```

Second Step: check how many instances per labels there are

```
In [12]: dataset["ocean_proximity"].value_counts()
```

```
Out[12]: <1H OCEAN      9034  
        INLAND        6496  
        NEAR OCEAN    2628  
        NEAR BAY      2270  
        ISLAND         5  
        Name: ocean_proximity, dtype: int64
```

Label *ISLAND* is removed since there are only 5 examples

```
In [13]: dataset[dataset["ocean_proximity"]=="ISLAND"]
```

```
Out[13]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
8314	-118.32	33.35	27.0	1675.0	521.0	
8315	-118.33	33.34	52.0	2359.0	591.0	

8316	-118.32	33.33	52.0	2127.0	512.0
8317	-118.32	33.34	52.0	996.0	264.0
8318	-118.48	33.43	29.0	716.0	214.0

	population	households	median_income	median_house_value	\
8314	744.0	331.0	2.1579	450000.0	
8315	1100.0	431.0	2.8333	414700.0	
8316	733.0	288.0	3.3906	300000.0	
8317	341.0	160.0	2.7361	450000.0	
8318	422.0	173.0	2.6042	287500.0	

	ocean_proximity
8314	ISLAND
8315	ISLAND
8316	ISLAND
8317	ISLAND
8318	ISLAND

```
In [14]: dataset.drop(index=range(8314,8319),inplace=True)
```

```
In [15]: dataset["ocean_proximity"].value_counts()
```

```
Out[15]: <1H OCEAN      9034
INLAND      6496
NEAR OCEAN    2628
NEAR BAY     2270
Name: ocean_proximity, dtype: int64
```

```
In [16]: labels=list(set(dataset["ocean_proximity"]))
labels
```

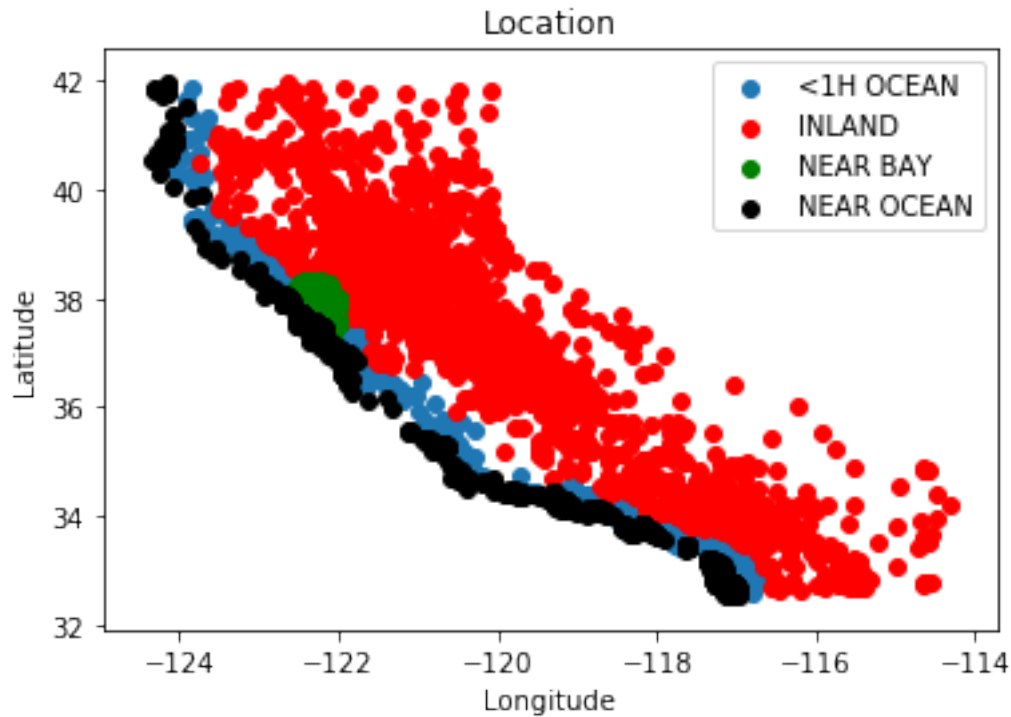
```
Out[16]: ['NEAR OCEAN', '<1H OCEAN', 'INLAND', 'NEAR BAY']
```

```
In [17]: dataset.shape
```

```
Out[17]: (20428, 10)
```

Location of the block groups (houses) included in the clean dataset:

```
In [18]: plt.title("Location")
plt.xlabel("Longitude")
plt.ylabel("Latitude")
plt.scatter(dataset[dataset["ocean_proximity"]=="<1H OCEAN"]["longitude"],
            dataset[dataset["ocean_proximity"]=="<1H OCEAN"]["latitude"])
plt.scatter(dataset[dataset["ocean_proximity"]=="INLAND"]["longitude"],
            dataset[dataset["ocean_proximity"]=="INLAND"]["latitude"],c="r")
plt.scatter(dataset[dataset["ocean_proximity"]=="NEAR BAY"]["longitude"],
            dataset[dataset["ocean_proximity"]=="NEAR BAY"]["latitude"],c="g")
plt.scatter(dataset[dataset["ocean_proximity"]=="NEAR OCEAN"]["longitude"],
            dataset[dataset["ocean_proximity"]=="NEAR OCEAN"]["latitude"],c="k")
plt.legend(["<1H OCEAN", "INLAND", "NEAR BAY", "NEAR OCEAN"])
plt.show()
```



Third Step: suffle the dataset (three times)

```
In [19]: dataset=dataset.sample(frac=1) #frac is the fraction of axis items to return. 1 means a
dataset=dataset.sample(frac=1)
dataset=dataset.sample(frac=1).reset_index(drop=True) #Reset index and drop the old one
dataset.head()
```

```
Out[19]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-124.05	40.85	31.0	2414.0	428.0	
1	-122.19	39.50	23.0	462.0	97.0	
2	-117.20	33.70	23.0	6323.0	1196.0	
3	-119.55	36.71	32.0	1963.0	508.0	
4	-119.22	34.18	17.0	3332.0	762.0	

	population	households	median_income	median_house_value	ocean_proximity
0	1005.0	401.0	3.5156	143000.0	NEAR OCEAN
1	261.0	90.0	2.1705	53000.0	INLAND
2	1984.0	1124.0	2.3276	92400.0	<1H OCEAN
3	2052.0	518.0	1.9076	55800.0	INLAND
4	1797.0	673.0	4.4292	231200.0	NEAR OCEAN

Fourth Step: split dataset vertically into attributes x and label t for supervised learning

```
In [20]: t = dataset["ocean_proximity"]
t[:10]
```

```
Out [20]: 0    NEAR OCEAN
          1      INLAND
          2    <1H OCEAN
          3      INLAND
          4    NEAR OCEAN
          5    <1H OCEAN
          6    NEAR OCEAN
          7    <1H OCEAN
          8    <1H OCEAN
          9      INLAND
          Name: ocean_proximity, dtype: object
```

```
In [21]: x = dataset.drop (columns="ocean_proximity")
```

```
In [22]: x.head()
```

```
Out [22]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-124.05	40.85	31.0	2414.0	428.0	
1	-122.19	39.50	23.0	462.0	97.0	
2	-117.20	33.70	23.0	6323.0	1196.0	
3	-119.55	36.71	32.0	1963.0	508.0	
4	-119.22	34.18	17.0	3332.0	762.0	

	population	households	median_income	median_house_value
0	1005.0	401.0	3.5156	143000.0
1	261.0	90.0	2.1705	53000.0
2	1984.0	1124.0	2.3276	92400.0
3	2052.0	518.0	1.9076	55800.0
4	1797.0	673.0	4.4292	231200.0

Fifth Step: discretizing and one-hot encoding of labels (target values)

```
In [23]: encoder = LabelEncoder() # Function that transform non-numeral labels into integers.
          discretized_t = encoder.fit_transform(t.values)
          labels = [encoder.inverse_transform([value]) for value in range(4)] #hold the label for
          discretized_t[:10]
```

```
Out [23]: array([3, 1, 0, 1, 3, 0, 3, 0, 0, 1])
```

```
In [24]: discretized_t = pd.DataFrame(data=discretized_t,columns=["ocean_proximity"])
          discretized_t[:10]
```

```
Out [24]:
```

	ocean_proximity
0	3
1	1
2	0
3	1
4	3
5	0

```

6          3
7          0
8          0
9          1

```

```
In [25]: labels = np.array(labels).reshape(1,-1)[0]
labels
```

```
Out[25]: array(['<1H OCEAN', 'INLAND', 'NEAR BAY', 'NEAR OCEAN'], dtype=object)
```

```
In [26]: { value: encoder.inverse_transform([value]) for value in range(4)}
```

```
Out[26]: {0: array(['<1H OCEAN'], dtype=object),
1: array(['INLAND'], dtype=object),
2: array(['NEAR BAY'], dtype=object),
3: array(['NEAR OCEAN'], dtype=object)}
```

```
In [27]: encoder = OneHotEncoder(sparse=False) # Function that one-hot encoders integers
#one_hot_t = encoder.fit_transform(discretized_t.reshape(-1,1)).toarray()
one_hot_t = encoder.fit_transform (t.values.reshape(-1,1))
one_hot_t[:10]
```

```
Out[27]: array([[0., 0., 0., 1.],
[0., 1., 0., 0.],
[1., 0., 0., 0.],
[0., 1., 0., 0.],
[0., 0., 0., 1.],
[1., 0., 0., 0.],
[0., 0., 0., 1.],
[1., 0., 0., 0.],
[1., 0., 0., 0.],
[0., 1., 0., 0.]])
```

```
In [28]: one_hot_t = pd.DataFrame(data=one_hot_t,columns=labels)
one_hot_t[:10]
```

```
Out[28]:
```

	<1H OCEAN	INLAND	NEAR BAY	NEAR OCEAN
0	0.0	0.0	0.0	1.0
1	0.0	1.0	0.0	0.0
2	1.0	0.0	0.0	0.0
3	0.0	1.0	0.0	0.0
4	0.0	0.0	0.0	1.0
5	1.0	0.0	0.0	0.0
6	0.0	0.0	0.0	1.0
7	1.0	0.0	0.0	0.0
8	1.0	0.0	0.0	0.0
9	0.0	1.0	0.0	0.0

Sixth Step: Normalization of the input dataset within the range [-1,1]


```
In [29]: x = pd.DataFrame (minmax_scale (x, feature_range=(-1,1)),columns=x.columns)
```

```
In [30]: x[:10]
```

```
Out [30]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-0.940239	0.766206	0.176471	-0.877308	-0.867474	
1	-0.569721	0.479277	-0.137255	-0.976601	-0.970205	
2	0.424303	-0.753454	-0.137255	-0.678468	-0.629112	
3	-0.043825	-0.113709	0.215686	-0.900249	-0.842644	
4	0.021912	-0.651435	-0.372549	-0.830612	-0.763811	
5	-0.679283	0.336876	-0.568627	-0.770436	-0.718187	
6	0.235060	-0.738576	0.019608	-0.845313	-0.756052	
7	0.217131	-0.647184	0.568627	-0.894959	-0.860025	
8	0.207171	-0.670563	1.000000	-0.935399	-0.900372	
9	-0.057769	-0.190223	-0.647059	-0.813215	-0.772502	

	population	households	median_income	median_house_value	
0	-0.943833	-0.868443	-0.584047	-0.472163	
1	-0.985538	-0.970728	-0.769576	-0.843295	
2	-0.888954	-0.630653	-0.747907	-0.680822	
3	-0.885143	-0.829962	-0.805837	-0.831749	
4	-0.899437	-0.778984	-0.458035	-0.108453	
5	-0.913058	-0.826673	-0.606461	-0.687420	
6	-0.921803	-0.754646	-0.578916	0.056082	
7	-0.928305	-0.850682	-0.299913	0.216082	
8	-0.945010	-0.884230	-0.806016	0.016907	
9	-0.895681	-0.764184	-0.708335	-0.730719	

Some descriptive statistics on the attributes to confirm the max and min values

```
In [31]: x.describe()
```

```
Out [31]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20428.000000	20428.000000	20428.000000	20428.000000	
mean	-0.048005	-0.342449	0.083519	-0.865977	
std	0.399150	0.454052	0.493732	0.111168	
min	-1.000000	-1.000000	-1.000000	-1.000000	
25%	-0.492032	-0.704570	-0.333333	-0.926344	
50%	0.165339	-0.634431	0.098039	-0.891907	
75%	0.262948	0.100956	0.411765	-0.840213	
max	1.000000	1.000000	1.000000	1.000000	

	total_bedrooms	population	households	median_income	\
count	20428.000000	20428.000000	20428.000000	20428.000000	
mean	-0.833365	-0.920282	-0.836051	-0.534967	
std	0.130796	0.063526	0.125745	0.261987	
min	-1.000000	-1.000000	-1.000000	-1.000000	
25%	-0.908442	-0.955997	-0.908239	-0.715383	
50%	-0.865301	-0.934808	-0.865812	-0.581026	

75%	-0.799503	-0.903585	-0.801677	-0.414605
max	1.000000	1.000000	1.000000	1.000000

	median_house_value
count	20428.000000
mean	-0.208981
std	0.475925
min	-1.000000
25%	-0.569173
50%	-0.320823
75%	0.029691
max	1.000000

The correlation matrix permits to visualize dependencies between pairs of attributes: values close to -1 or +1 indicate high correlation. A negative correlation value means that when the value of an attribute gets high the value of the other attribute decreases, and viceversa. Positive correlation values means that both attributes increase or decrease simultaneously. For example *population* and *households* have high positive correlation, while high negative correlations do not occur. *population* and *median_income* are highly uncorrelated.

In [32]: `x.corr()`

```
Out [32]:
```

	longitude	latitude	housing_median_age	total_rooms	\
longitude	1.000000	-0.924628	-0.109556	0.045555	
latitude	-0.924628	1.000000	0.012192	-0.036799	
housing_median_age	-0.109556	0.012192	1.000000	-0.360634	
total_rooms	0.045555	-0.036799	-0.360634	1.000000	
total_bedrooms	0.069653	-0.067066	-0.320486	0.930382	
population	0.100380	-0.109193	-0.295715	0.857273	
households	0.056604	-0.071939	-0.302714	0.918987	
median_income	-0.015464	-0.079796	-0.118191	0.197822	
median_house_value	-0.045642	-0.144312	0.106077	0.133516	

	total_bedrooms	population	households	median_income	\
longitude	0.069653	0.100380	0.056604	-0.015464	
latitude	-0.067066	-0.109193	-0.071939	-0.079796	
housing_median_age	-0.320486	-0.295715	-0.302714	-0.118191	
total_rooms	0.930382	0.857273	0.918987	0.197822	
total_bedrooms	1.000000	0.877758	0.979740	-0.007767	
population	0.877758	1.000000	0.907177	0.004989	
households	0.979740	0.907177	1.000000	0.013350	
median_income	-0.007767	0.004989	0.013350	1.000000	
median_house_value	0.049792	-0.025069	0.065122	0.688848	

	median_house_value
longitude	-0.045642
latitude	-0.144312
housing_median_age	0.106077

total_rooms	0.133516
total_bedrooms	0.049792
population	-0.025069
households	0.065122
median_income	0.688848
median_house_value	1.000000

Finally, both attributes matrix x and target labels t are saved to csv files.

```
In [33]: x.to_csv(ATT_FILE_NAME, index=False)
         #discretized_t.to_csv(DISCRETIZED_LABEL_FILE_NAME, index=False)
         one_hot_t.to_csv(ONE_HOT_ENCODED_LABEL_FILE_NAME, index=False)
```